

# □ 1. Sentiment Analysis of Product Reviews

**Goal:** Predict if a review is *positive* or *negative*.

**What you'll do:**

- In SageMaker Studio, open the “JumpStart” tab.
- Choose the “**Text Classification (Sentiment Analysis)**” model (pre-built by AWS).
- Use a small sample CSV file (e.g. 10–20 Amazon review lines).
- Click “**Deploy**” — SageMaker automatically sets up the model endpoint.
- Type a few example sentences and see if the model predicts *positive* or *negative*.

**Concept learned:** Natural Language Processing (NLP) and real-time inference.

---

# □ 2. Image Classification — Identify Cats vs. Dogs

**Goal:** Use a pretrained image model to classify images.

**Steps:**

- Go to **SageMaker JumpStart** → **Image Classification**.
- Select a pre-trained model (e.g. ResNet-50).
- Upload 5–10 images of cats and dogs.
- Deploy the model — SageMaker handles the rest.
- Test each image and see which label it predicts.

**Concept learned:** Computer Vision and image inference.

---

# □ 3. Predict House Prices (Simple Regression)

**Goal:** Estimate house prices based on a few features like size, bedrooms, location.

**Steps:**

- Download a small dataset (e.g. from Kaggle or a CSV of 20 rows).
- Open SageMaker Studio → **AutoML (Autopilot)**.

- Upload your CSV and select “**Predict price**” as target.
- SageMaker automatically builds, trains, and evaluates a model.
- View the model report and make a few test predictions.

**Concept learned:** AutoML and supervised regression.

---

## 4. Sales Forecasting

**Goal:** Predict next week’s sales from past sales data.

**Steps:**

- Prepare a 2-column CSV: `date`, `sales`.
- Use **JumpStart** → **Time Series Forecasting**.
- Upload the CSV, deploy the pretrained model.
- Check the predicted sales values and trend chart.

**Concept learned:** Forecasting and time series prediction.

---

## 5. Detect Fake News Headlines

**Goal:** Classify whether a headline is “Real” or “Fake”.

**Steps:**

- Use **JumpStart** → **Text Classification (Binary)**.
- Upload a small dataset with two columns: `headline`, `label` (“real”/“fake”).
- Deploy the pre-trained model.
- Type a few headlines manually to test.

**Concept learned:** Text classification and binary prediction.

---

## Tips to Stay Free-Tier Safe

- Always pick **ml.t3.medium** or **ml.t2.medium** instance types (free-tier eligible).
  - **Stop or delete** endpoints once you’re done (otherwise you’ll be billed).
  - Use **JumpStart** or **Autopilot** instead of writing Python code — these are *click-and-run* tools.
-